

Caret / Recursive Partitioning

Saqib Ali

May 16th, 2017

Load the data

```
data.joined <- readRDS(file="/home/saqib/ml_at_berkeley/CSX460/04-logistic-regression/04-exercise-nycfl
```

Add a categorical variable for `arr_delay >= 22` minutes. It is called `arrival_delayed`

```
data.joined$arrival_delayed <- factor(ifelse(data.joined$arr_delay >= 22, 1,0))
```

Filter out rows with NAs

```
data.joined <- data.joined %>% filter(!is.na(arrival_delayed))
```

```
#data.joined <- data.joined[, speed:=NULL]
```

```
data.joined.training <- sample_frac(data.joined, .75)
```

```
data.joined.testing <- sample_frac(data.joined, .5)
```

Exercise 1: caret/logistic regression (5 points)

Rebuild your logistic regression model from the previous week, this time using the `caret` package.

- Calculate the training or apparent performance of the model.
- Calculate an unbiased measure of performance
- Create a ROC Curve for your model

Show all work.

```
# Your Work Here
```

```
#data.joined <- data.joined[, speed:=NULL]
```

```
#lapply(data.joined, levels)
```

```
##(l <- sapply(data.joined, function(x) is.factor(x)))
```

```
fitControl <- trainControl(method = "cv", number = 2)
```

```
glmFit <- train(arrival_delayed ~ dep_delay + dest + origin + year + month + day + hour + sched_dep_time
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =  
## ifelse(type == : prediction from a rank-deficient fit may be misleading
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =  
## ifelse(type == : prediction from a rank-deficient fit may be misleading
```

```

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
pred <- as.vector(ifelse(predict(glmFit, newdata=data.joined, type="prob"), "1" < .5, 0, 1))

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading
probsTest <- predict(glmFit, data.joined.testing, type = "prob")

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading
pred <- factor( ifelse(probsTest[, "1"] > 0.5, "1", "0") )
confusionMatrix(pred, data.joined.testing$arrival_delayed)


## Confusion Matrix and Statistics
##
##              Reference
## Prediction      0      1
##              0 93240  6579
##              1  1994 15009
##
##              Accuracy : 0.9266
##              95% CI : (0.9251, 0.9281)
##              No Information Rate : 0.8152
##              P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 0.7346
##              Mcnemar's Test P-Value : < 2.2e-16
##
##              Sensitivity : 0.9791
##              Specificity : 0.6952
##              Pos Pred Value : 0.9341
##              Neg Pred Value : 0.8827
##              Prevalence : 0.8152
##              Detection Rate : 0.7981
##              Detection Prevalence : 0.8545
##              Balanced Accuracy : 0.8372
##
##              'Positive' Class : 0
##
library(pROC)

## Type 'citation("pROC")' for a citation.
##
## Attaching package: 'pROC'
##
## The following objects are masked from 'package:stats':
##
##      cov, smooth, var
probsTrain <- predict(glmFit, data.joined.training, type = "prob")

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading

```

```
rocCurve <- roc(response = data.joined.training$arrival_delayed,
  predictor = probsTrain[, "1"],
  levels = rev(levels(data.joined.training$arrival_delayed)))
plot(rocCurve, print.thres = "best")
```



Exercise 2: caret/rpart (5 points)

Using the `caret` and `rpart` packages, create a **classification** model for flight delays using your NYC FLight data. Your solution should include:

- The use of `caret` and `rpart` to train a model.
- An articulation of the the problem your are
- An naive model
- An unbiased calculation of the performance metric
- A plot of your model – (the actual tree; there are several ways to do this)
- A discussion of your model

Show and describe all work

```
# Your Work Here
```

Questions:

- Discuss the difference between the models and why you would use one model over the other?
- How might you produce an ROC type curve for the `rpart` model?